

# Vector Space Classification

Cornelia Caragea

Acknowledgments: Manning, Rai

Department of Computer Science  
University of Illinois at Chicago


# Required Reading

- ▶ “Information Retrieval” textbook
  - ▶ Chapter 14: Vector Space Classification


# Classification Using Vector Spaces

- ▶ In vector space classification, the training set corresponds to a labeled set of points (equivalently, vectors)
- ▶ **Premise 1:** Documents in the same class form a contiguous region of space
- ▶ **Premise 2:** Documents from different classes do not overlap (much)
- ▶ Learning a classifier: build surfaces to delineate classes in the space


# Documents in a Vector Space



# Test Document of What Class?



# Test Document = Government



Our focus: how to find “good” separators (or decision boundaries)?

# Two Vector Space Classification Methods

- ▶ Rocchio Classification
  - ▶ Divides the vector space into regions centered on **centroids**
  - ▶ Simple and efficient, but inaccurate if classes are not approximately spheres with similar radii.
- ▶ K-Nearest Neighbors Classification
  - ▶ Assigns the majority class of the k nearest neighbors to a test document
  - ▶ Requires no explicit training
  - ▶ It is less efficient than other classification methods

# Rocchio Classification

- ▶ Uses **centroids** to define the boundaries
- ▶ The **centroid of a class  $c$**  is computed as the vector average:


$$\mu(c) = \frac{1}{|D_c|} \sum_{d \in D_c} v(d)$$

where  $D_c$  is the set of documents in  $D$  whose class is  $c$  and  $v(d)$  is the vector space representation of  $d$ .

- ▶ The **boundary between two classes** in Rocchio classification is the set of points with equal distance from the two centroids.



# Example



# Rocchio Classification: Training and Testing

TRAINROCCHIO( $\mathbb{C}, \mathbb{D}$ )

```
1  for each  $c_j \in \mathbb{C}$   
2  do  $D_j \leftarrow \{d : \langle d, c_j \rangle \in \mathbb{D}\}$   
3       $\vec{\mu}_j \leftarrow \frac{1}{|D_j|} \sum_{d \in D_j} \vec{v}(d)$   
4  return  $\{\vec{\mu}_1, \dots, \vec{\mu}_J\}$ 
```

APPLYROCCHIO( $\{\vec{\mu}_1, \dots, \vec{\mu}_J\}, d$ )

```
1  return  $\arg \min_j |\vec{\mu}_j - \vec{v}(d)|$ 
```

The distance to a centroid is computed as the Euclidian distance.

# Rocchio Classification


- ▶ Rocchio forms a simple representative for each class: the centroid/prototype
- ▶ Classification: nearest prototype/centroid
- ▶ Cheap to train and test documents

# k Nearest Neighbor Classification

- ▶  $kNN = k$  Nearest Neighbor
- ▶ To classify a document  $d$ :
  - ▶ Define  $k$ -neighborhood as the  $k$  nearest neighbors of  $d$
  - ▶ Pick the majority class label in the  $k$ -neighborhood

## $K$ Nearest Neighbor ( $K$ -NN)

- ▶ Given training  $\mathcal{D} = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and a test point
- ▶ Prediction Rule: Look at the  $K$  most similar training examples



- ▶ Assign the majority class label (**majority voting**)
- ▶ The algorithm requires:
  - ▶ Parameter  $K$ : number of nearest neighbors to look for
  - ▶ **Distance function**: To compute the similarities between examples
- ▶ **Special Case**: 1-Nearest Neighbor

# $K$ Nearest Neighbors Algorithm

- ▶ Compute the test point's distance from each training point
- ▶ Sort the distances in ascending (or descending) order
- ▶ Use the sorted distances to select the  $K$  nearest neighbors
- ▶ Use **majority rule** (for classification)

**Note:**  $K$ -Nearest Neighbors is called a *non-parametric* method

- ▶ Unlike other supervised learning algorithms,  $K$ -Nearest Neighbors does not learn an explicit mapping  $f$  from the training data
- ▶ It simply uses the training data at test time to make predictions

# K-NN: Feature Normalization

- ▶ Note: Features should be on the same scale
- ▶ Example: if one feature has its values in millimeters and another has in centimeters, we would need to normalize
- ▶ One way is:
  - ▶ Replace  $x_{im}$  by  $z_{im} = \frac{(x_{im} - \bar{x}_m)}{\sigma_m}$  (make them zero mean, unit variance)
  - ▶  $\bar{x}_m = \frac{1}{N} \sum_{i=1}^N x_{im}$ : empirical mean of  $m^{th}$  feature
  - ▶  $\sigma_m^2 = \frac{1}{N} \sum_{i=1}^N (x_{im} - \bar{x}_m)^2$ : empirical variance of  $m^{th}$  feature

# K-NN: Computing the Distances

- ▶ The  $K$ -NN algorithm requires computing distances of the test example from each of the training examples
- ▶ Several ways to compute distances
- ▶ The choice depends on the **type of the features** in the data
- ▶ Real-valued features ( $x_i \in R^D$ ): **Euclidean distance** is commonly used

$$d(x_i, x_j) = \sqrt{\sum_{m=1}^D (x_{im} - x_{jm})^2} = \sqrt{\|x_i\|^2 + \|x_j\|^2 - 2x_i^T x_j}$$



## K-NN: Some Other Distance Measures

- ▶ Binary-valued features
  - ▶ Use Hamming distance:  $d(x_i, x_j) = \sum_{m=1}^D I(x_{im} \neq x_{jm})$
  - ▶ Hamming distance counts the number of features where the two examples disagree
- ▶ Mixed feature types (some real-valued and some binary-valued)?
  - ▶ Can use mixed distance measures
  - ▶ E.g., Euclidean for the real part, Hamming for the binary part
- ▶ Can also assign weights to features:  
$$d(x_i, x_j) = \sum_{m=1}^D w_m d(x_{im}, x_{jm})$$

# k Nearest Neighbor

- ▶ Using only the closest example (1NN) subject to errors due to:
  - ▶ A single atypical example.
  - ▶ Noise (i.e., an error) in the category label of a single training example.
- ▶ More robust: find the  $k$  examples and return the majority category of these  $k$
- ▶  $k$  is typically odd to avoid ties; 3 and 5 are most common
- ▶ Choosing  $k$ 
  - ▶ Often data dependent and heuristic based
  - ▶ Or using [cross-validation](#) (using some held-out data)
  - ▶ In general, a  $k$  too small or too big is bad!

# kNN: Discussion

- ▶ No training necessary
- ▶ Scales well with large number of classes
  - ▶ No need to train  $n$  classifiers for  $n$  classes
- ▶ It could be more accurate than NB or Rocchio

# K-Nearest Neighbor: Further Discussion

- ▶ What is nice
  - ▶ Simple and intuitive; easily implementable
- ▶ What is not so nice...
  - ▶ Store all the training data in memory even at test time
    - ▶ Can be memory intensive for large training datasets
    - ▶ An example of **non-parametric**, or memory/instance-based methods
    - ▶ Different from **parametric, model-based** learning models
  - ▶ Expensive at test time:  $O(ND)$  computations for each test point
    - ▶ Have to **search through all training data** to find nearest neighbors
    - ▶ Distance computations with  $N$  training points ( $D$  features each)

# Which classifier do I use for a given text classification problem?

- ▶ Is there a learning method that is optimal for all text classification problems?
  - ▶ No
- ▶ Factors to take into account:
  - ▶ How much training data is available?
  - ▶ How simple/complex is the problem? (linear vs. nonlinear decision boundary)
  - ▶ How noisy is the data?
  - ▶ How stable is the problem over time?