

Information Retrieval and Web Search

Cornelia Caragea

Computer Science
University of Illinois at Chicago

Credits for slides: Mooney

Retrieval Models

Required Reading

- “Information Retrieval” textbook
 - Chapter 1: Boolean retrieval
 - Chapter 6: Scoring, term weighting & the vector space model

Classes of Retrieval Models

- Exact match
 - Boolean models (set theoretic)
- Ranking - “Best” match
 - Vector space models (algebraic)
 - Probabilistic models

Exact vs. Best Match

- Exact-match
 - Query specifies precise retrieval criteria
 - Every document either matches or fails to match query
 - Result is a **set** of documents
 - **Unordered** in pure exact match
- Best-match
 - Query describes good or “best” matching document
 - Every document matches query to some degree
 - Result is a **ranked list** of documents

Exact-match Pros & Cons

- Advantages of exact match
 - Can be very efficiently implemented
 - Predictable, easy to explain
 - Structured queries for pinpointing precise documents - very expressive
 - Works well when you know exactly (or roughly) what the collection contains and what you are looking for
- Disadvantages of exact match
 - Query formulation difficult for most users
 - Difficulty increases with collection size
 - The indexing vocabulary must be the same as query vocabulary
 - Ranking models are consistently shown to be better


Boolean Retrieval Models

- Exact-match models
 - Simple models based on set theory
 - Neat formalism, precise semantic $q = a \wedge (b \vee \neg c)$
 - Queries are logic expressions with document features as operands, and specify precise relevance criteria
 - The models retrieve documents iff they satisfy a **Boolean expression**
 - Documents are returned in no particular order
- Supported operators (query language)
 - **Logical operators**: AND, OR, NOT
 - Most systems support simple regular expressions as search terms to match spelling variants
 - colou?r
 - ab*c
 - ab+c

Boolean Model

- Consider

- $q = a \wedge (b \vee \neg c)$




Boolean Model

- Consider

- $q = a \wedge (b \vee \neg c)$

- Result: $(1, 1, 1) \vee (1, 1, 0) \vee (1, 0, 0)$



Drawbacks of the Boolean Model


- Retrieval is based on binary decision criteria with no notion of partial matching
- No ranking of the documents is provided (absence of a grading scale)
- Information need has to be translated into a Boolean expression which most users find awkward
- The Boolean queries formulated by the users are most often too simplistic

Best-Match Retrieval

- Best-match or ranking models are more common
- Advantages:
 - Significantly more effective than exact match
 - Easier to use (supports full text queries)
- Disadvantages:
 - Efficiency is always less than exact match (cannot reject documents early)
- Boolean or structured queries can be part of a best-match retrieval model

Vector Space Models

- **Key idea:** Everything (documents, queries, terms) is a vector in a high-dimensional space.




- The geometry of space induces a **similarity measure** between documents
- The documents are ranked based on their similarity with the query
- History:
 - Invented by Gerald Salton (1960/70)

Issues for Vector Space Models

- How to determine important words in a document?
 - How to select basis vectors (dimensions)
- How to convert objects into vectors?
 - Documents, queries, terms
- Assumption - not all terms are equally useful for representing the document contents, less frequent terms allow identifying a narrower set of documents
 - The importance of the index terms is represented by weights associated to them.
 - How to determine the degree of importance of a term within a document and within the entire collection?
- How to compare objects in the vector space?
 - How to determine the degree of similarity between a document and the query?
- In the case of the web, what is a collection?

Example Graphical Representation

- $D_1 = (2T_1, 3T_2, 5T_3)$
- $D_2 = (3T_1, 7T_2, 1T_3)$
- $Q = (0T_1, 0T_2, 2T_3)$



- Is D_1 or D_2 more similar to Q ?
- How to measure the degree of similarity? Distance? Angle?

The Vector-Space Model

- Assume t distinct terms remain after preprocessing; call them **index terms** or the **vocabulary**.
- These “orthogonal” terms form a basis of a vector space.
Dimension $= t = |\text{vocabulary}|$
- Each term, i , in a document or query, j , is given a real-valued weight, w_{ij} .
- Both documents and queries are expressed as t -dimensional vectors:
$$d_j = (w_{1j}, w_{2j}, \dots, w_{tj})$$

Document Collection

- A collection of n documents can be represented in the vector space model by a term-document matrix.
- An entry in the matrix corresponds to the “weight” of a term in the document; zero means the term has no significance in the document or it simply does not exist in the document.

$$\begin{pmatrix} & T_1 & T_2 & \dots & T_t \\ D_1 & w_{11} & w_{21} & \dots & w_{t1} \\ D_2 & w_{12} & w_{22} & \dots & w_{t2} \\ \vdots & \vdots & \vdots & & \vdots \\ \vdots & \vdots & \vdots & & \vdots \\ D_n & w_{1n} & w_{2n} & \dots & w_{tn} \end{pmatrix}$$

Term Weights: Term Frequency

- More frequent terms in a document are more important, i.e. more indicative of the topic.
 f_{ij} = frequency of term i in document j
- May want to normalize *term frequency* (*tf*)
 - e.g. by dividing by the frequency of the most common term in the document:

$$tf_{ij} = \frac{f_{ij}}{\max_i \{f_{ij}\}}$$

Term Weights: Inverse Document Frequency

- Terms that appear in many *different* documents are less indicative of the overall topic.
 - df_i = document frequency of term i = number of documents containing term i
 - idf_i = inverse document frequency of term $i = \log_2(N/df_i)$ (N : total number of documents)
- An indication of a term's *discrimination* power.
- Log used to dampen the effect relative to tf .

TF-IDF Weighting

- A typical combined term importance indicator is *tf-idf* weighting:

$$w_{ij} = tf_{ij}idf_i = tf_{ij} \log_2(N/df_i)$$

- A term occurring frequently in the document but rarely in the rest of the collection is given high weight.
- Experimentally, *tf-idf* has been found to work well.

Computing *tf-idf* - An Example

- Given a document containing terms with given frequencies:

$$A(3), B(2), C(1)$$

- Assume collection contains 10,000 documents and document frequencies of these terms are:

$$A(50), B(1300), C(250)$$

- Compute *tf*, *idf*, *tf-idf*?

$$w_{ij} = tf_{ij}idf_i = (f_{ij}/\max_i\{f_{ij}\}) \cdot \log_2(N/df_i)$$

Computing *tf-idf* - An Example

- Given a document containing terms with given frequencies:

$$A(3), B(2), C(1)$$

- Assume collection contains 10,000 documents and document frequencies of these terms are:

$$A(50), B(1300), C(250)$$

- Then:

$$A : tf = 3/3; idf = \log_2(10000/50) = 7.6; tf-idf = 7.6$$

$$B : tf = 2/3; idf = \log_2(10000/1300) = 2.9; tf-idf = 2.0$$

$$C : tf = 1/3; idf = \log_2(10000/250) = 5.3; tf-idf = 1.8$$

Query Vector

- Query vector is typically treated as a document and is also *tf-idf* weighted.
- The alternative is for the user to supply weights for the given query terms.
 - Weighted query terms:
 $Q = \langle \text{database } 0.5; \text{ text } 0.8; \text{ information } 0.2 \rangle$
 - Unweighted query terms:
 $Q = \langle \text{database}; \text{ text}; \text{ information} \rangle$

Similarity Measure

- A **similarity measure** is a function that computes the **degree of similarity** between two vectors.
- Using a similarity measure between the query and each document:
 - It is possible to rank the retrieved documents in the order of presumed relevance.
 - It is possible to enforce a certain threshold so that the size of the retrieved set can be controlled.

Desiderata for Proximity

- If d_1 is near d_2 , then d_2 is near d_1 .
- If d_1 near d_2 , and d_2 near d_3 , then d_1 is not far from d_3 .
- No document is closer to d than d itself.

Vector Space Similarity: Common Measures

<u>Sim(X,Y)</u>	<u>Binary Term Vectors</u>	<u>Weighted Term Vectors</u>
Inner product	$ X \cap Y $	$\sum x_i \cdot y_i$
Dice coefficient	$\frac{2 X \cap Y }{ X + Y }$	$\frac{2\sum x_i \cdot y_i}{\sum x_i^2 + \sum y_i^2}$
Cosine coefficient	$\frac{ X \cap Y }{\sqrt{ X } \sqrt{ Y }}$	$\frac{\sum x_i \cdot y_i}{\sqrt{\sum x_i^2} \cdot \sqrt{\sum y_i^2}}$
Jaccard coefficient	$\frac{ X \cap Y }{ X + Y - X \cap Y }$	$\frac{\sum x_i \cdot y_i}{\sum x_i^2 + \sum y_i^2 - \sum x_i \cdot y_i}$

Inner Product

- Similarity between vectors for the document d_j and query q can be computed as the vector inner product (or the dot product):

$$\text{sim}(d_j, q) = d_j \cdot q = \sum_{i=1}^t w_{ij} w_{iq}$$

where w_{ij} is the weight of term i in document j and w_{iq} is the weight of term i in the query

- For binary vectors, the inner product is the number of matched query terms in the document (size of intersection).
- For weighted term vectors, it is the sum of the products of the weights of the matched terms.

Inner Product - Examples

- Binary:

	retrieval	database	architecture	computer	text	management	information
▪ D =	1	1	1	0	1	1	0
▪ Q =	1	0	1	0	0	1	1

Size of vector = size of vocabulary = 7;

0 means corresponding term not found in document or query

$$\text{sim}(D, Q) = ?$$

- Weighted:

$$D_1 = (2T_1, 3T_2, 5T_3), D_2 = (3T_1, 7T_2, 1T_3),$$

$$Q = (0T_1, 0T_2, 2T_3)$$

$$\text{sim}(D_1, Q) = ?$$

$$\text{sim}(D_2, Q) = ?$$

Inner Product - Examples

- Binary:

Size of vector = size of vocabulary
= 7; 0 means corresponding term
not found in document or query

$$\text{sim}(D, Q) = 3$$

	retrieval	database	architecture	computer	text	management	information
▪ D =	1	1	1	0	1	1	0
▪ Q =	1	0	1	0	0	1	1


- Weighted:

$$D_1 = (2T_1, 3T_2, 5T_3), D_2 = (3T_1, 7T_2, 1T_3),$$

$$Q = (0T_1, 0T_2, 2T_3)$$

$$\text{sim}(D_1, Q) = 2 \cdot 0 + 3 \cdot 0 + 5 \cdot 2 = 10$$

$$\text{sim}(D_2, Q) = 3 \cdot 0 + 7 \cdot 0 + 1 \cdot 2 = 2$$



Cosine Similarity Measure

- Cosine similarity measures the cosine of the angle between two vectors.
- Inner product normalized by the vector lengths.

$$\text{CosSim}(d_j, q) = \frac{\langle d_j, q \rangle}{\|d_j\| \cdot \|q\|} = \frac{\sum_{i=1}^t w_{ij} w_{iq}}{\sqrt{\sum_{i=1}^t w_{ij}^2 \cdot \sum_{i=1}^t w_{iq}^2}}$$

$$D_1 = (2T_1, 3T_2, 5T_3), D_2 = (3T_1, 7T_2, 1T_3),$$

$$Q = (0T_1, 0T_2, 2T_3)$$

$$\text{CosSim}(D_1, Q) = ?$$

$$\text{CosSim}(D_2, Q) = ?$$

Cosine Similarity Measure

- Cosine similarity measures the cosine of the angle between two vectors.
- Inner product normalized by the vector lengths.

$$\text{CosSim}(d_j, q) = \frac{\langle d_j, q \rangle}{\|d_j\| \cdot \|q\|} = \frac{\sum_{i=1}^t w_{ij} w_{iq}}{\sqrt{\sum_{i=1}^t w_{ij}^2 \cdot \sum_{i=1}^t w_{iq}^2}}$$

$$D_1 = (2T_1, 3T_2, 5T_3), D_2 = (3T_1, 7T_2, 1T_3),$$

$$Q = (0T_1, 0T_2, 2T_3)$$

$$\text{CosSim}(D_1, Q) = 10 / \sqrt{(4 + 9 + 25)(0 + 0 + 4)} = 0.81$$

$$\text{CosSim}(D_2, Q) = 2 / \sqrt{(9 + 49 + 1)(0 + 0 + 4)} = 0.13$$

D_1 is 6 times better than D_2 using cosine similarity but only 5 times better using inner product.

Vector Space Summary

- Very simple
 - Map everything to a vector
 - Compare using angle between vectors
- Challenge is mostly finding good weighting scheme
 - Variants on *tf-idf* are most common
- Another challenge is comparison function
 - Cosine comparison is most common
 - Generic inner product (without unit vectors) also occurs
- Considers both local (tf) and global (idf) word occurrence frequencies.
- Provides partial matching and ranked results.
- Tends to work quite well in practice despite obvious weaknesses.

Problems with Vector Space Model

- Missing semantic information (e.g. word sense).
- Missing syntactic information (e.g. phrase structure, word order, proximity information).
- Lacks the control of a Boolean model (e.g., requiring a term to appear in a document).
 - Given a two-term query “A B”, may prefer a document containing A frequently but not B, over a document that contains both A and B, but both less frequently
- Implementation?

Naïve Implementation

- Convert all documents in collection \mathcal{D} to *tf-idf* weighted vectors d_j for keyword vocabulary V .
- Convert query to a *tf-idf*-weighted vector q .
- For each d_j in \mathcal{D} do
 - Compute score $s_j = \text{CosSim}(d_j, q)$
- Sort documents by decreasing score.
- Present top ranked documents to the user.

Time complexity?

Naïve Implementation

- Convert all documents in collection \mathcal{D} to tf-idf weighted vectors d_j for keyword vocabulary V .
- Convert query to a tf-idf-weighted vector q .
- For each d_j in \mathcal{D} do
 - Compute score $s_j = \text{CosSim}(d_j, q)$
- Sort documents by decreasing score.
- Present top ranked documents to the user.

Time complexity: $O(|V| \cdot |D|)$ Bad for large V & D !

$|V| = 10,000$; $|D| = 100,000$; $|V| \cdot |D| = 1,000,000,000$

Practical Implementation

- Based on the observation that documents containing none of the query keywords do not affect the final ranking
- Try to identify only those documents that contain at least one query keyword
- Actual implementation of an inverted index