

CS 412 Introduction to Machine Learning

Logistic Regression – Code Tutorial

Instructor: Wei Tang

Department of Computer Science
University of Illinois at Chicago
Chicago IL 60607

<https://tangw.people.uic.edu>
tangw@uic.edu

Slides credit: Martín Pellarolo

Data

```
def generate_random_points(size=10, low=0, high=1):  
    data = (high - low) * np.random.random_sample((size, 2)) + low  
    return data
```

```
N = 20 # number of samples in each class
```

```
X1 = generate_random_points(N, 0, 1)  
y1 = np.ones(N)
```

```
X2 = generate_random_points(N, 1, 2)  
y2 = np.zeros(N)
```

```
X = np.concatenate((X1, X2), axis=0)  
y = np.concatenate((y1, y2), axis=0)
```

numpy.random.random_sample

`random.random_sample(size=None)`

Return random floats in the half-open interval [0.0, 1.0).

Results are from the “continuous uniform” distribution over the stated interval. To sample $Unif[a, b)$, $b > a$ multiply the output of `random_sample` by $(b-a)$ and add a :

```
(b - a) * random_sample() + a
```

Note

New code should use the `random` method of a `default_rng()` instance instead; please see the [Quick Start](#).

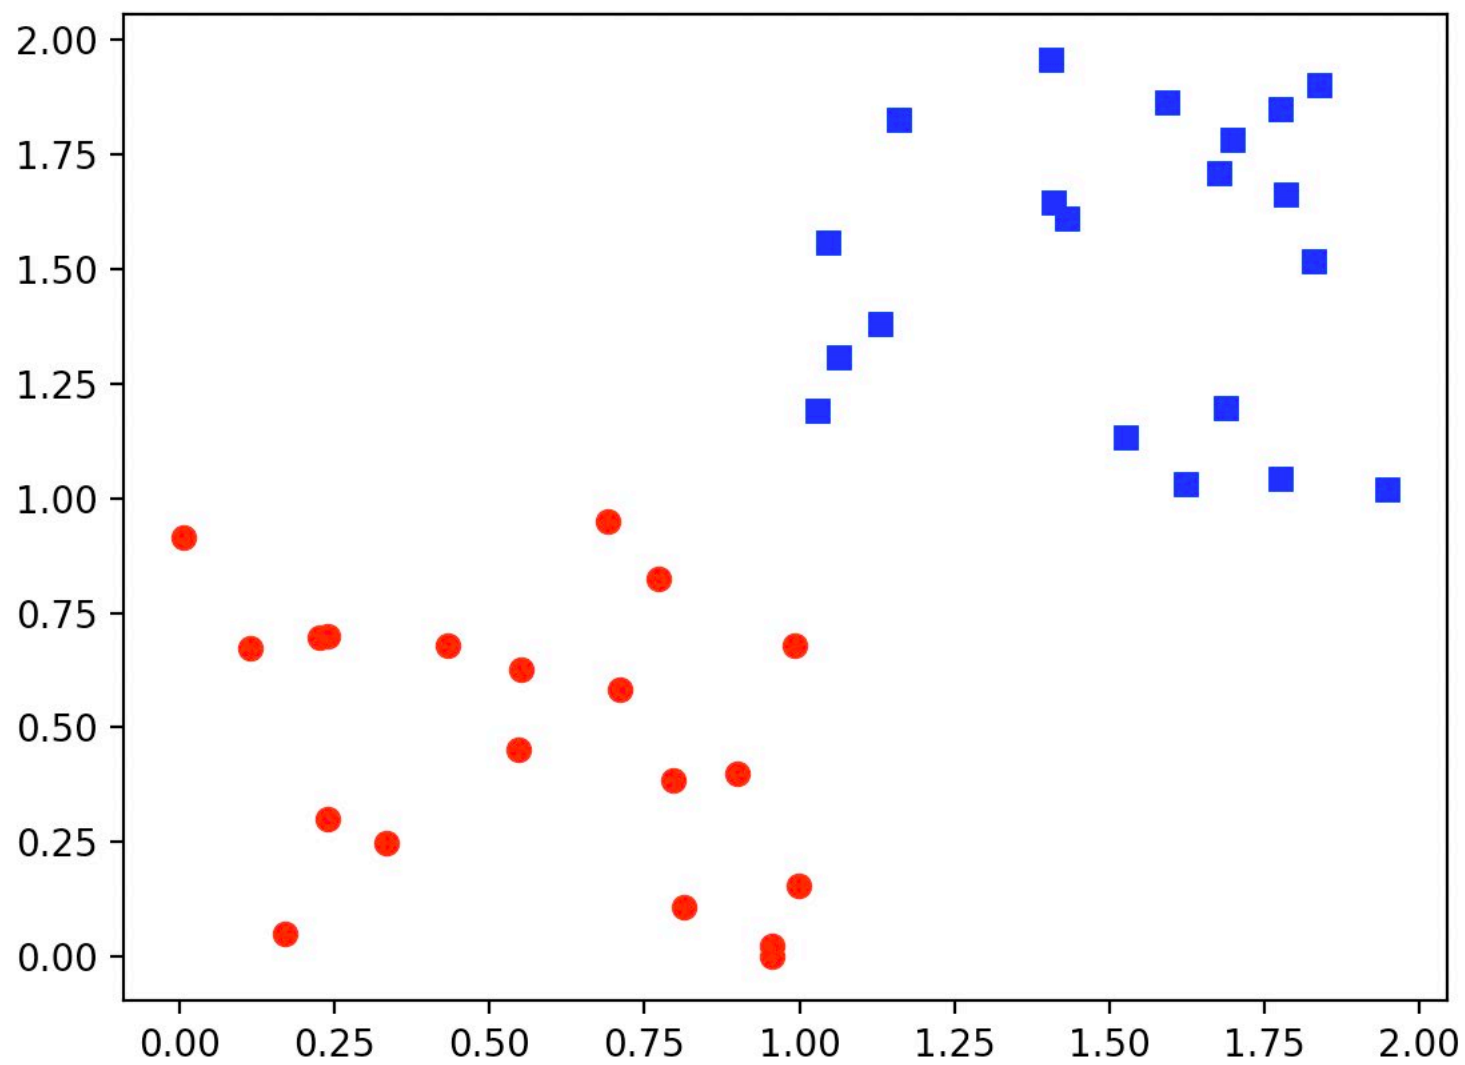
Parameters: `size` : *int or tuple of ints, optional*

Output shape. If the given shape is, e.g., (m, n, k) , then $m * n * k$ samples are drawn. Default is None, in which case a single value is returned.

Returns: `out` : *float or ndarray of floats*

Array of random floats of shape `size` (unless `size=None`, in which case a single float is returned).

Figure 1



x=0.887 y=1.866

Some useful functions

```
def sigmoid(z):  
    res = 1 / (1 + np.exp(-z))  
    return res
```

```
def add_intercept(X):  
    intercept = np.ones((X.shape[0], 1))  
    res = np.concatenate((intercept, X), axis=1)  
    return res
```

```
def loss(h, y):  
    res = (-y * np.log(h) - (1 - y) * np.log(1 - h)).mean()  
    return res
```

Learning

```
def learning(X, y, lr=0.01, num_iter=100000):  
    X = add_intercept(X)  
  
    # weights initialization  
    theta = np.zeros(X.shape[1])  
  
    for i in range(num_iter):  
        z = np.dot(X, theta)  
        h = sigmoid(z)  
        gradient = np.dot(X.T, (h - y)) / y.size  
        theta -= lr * gradient  
  
        if (i % 10000 == 0):  
            z = np.dot(X, theta)  
            h = sigmoid(z)  
            print(f'loss: {loss(h, y)} \t')  
    return theta
```

Prediction

```
def predict(X, theta, threshold=0.5):  
    X = add_intercept(X)  
    prob = sigmoid(np.dot(X, theta))  
    res = prob >= threshold  
    return res
```

The complete pipeline

```
N = 20 # number of samples in each class
```

```
X1 = generate_random_points(N, 0, 1)
```

```
y1 = np.ones(N)
```

```
X2 = generate_random_points(N, 1, 2)
```

```
y2 = np.zeros(N)
```

```
X = np.concatenate((X1, X2), axis=0)
```

```
y = np.concatenate((y1, y2), axis=0)
```

```
theta = learning(X, y)
```

```
y_preds = predict(X, theta)
```

```
print("Average classification accuracy:", (y_preds == y).mean())
```

```
plt.plot(X1[:,0], X1[:,1], 'ro', X2[:,0], X2[:,1], 'bs')
```

```
#  $x_2 = a \cdot x_1 + b$ 
```

```
a = -theta[1]/theta[2]
```

```
b = -theta[0]/theta[2]
```

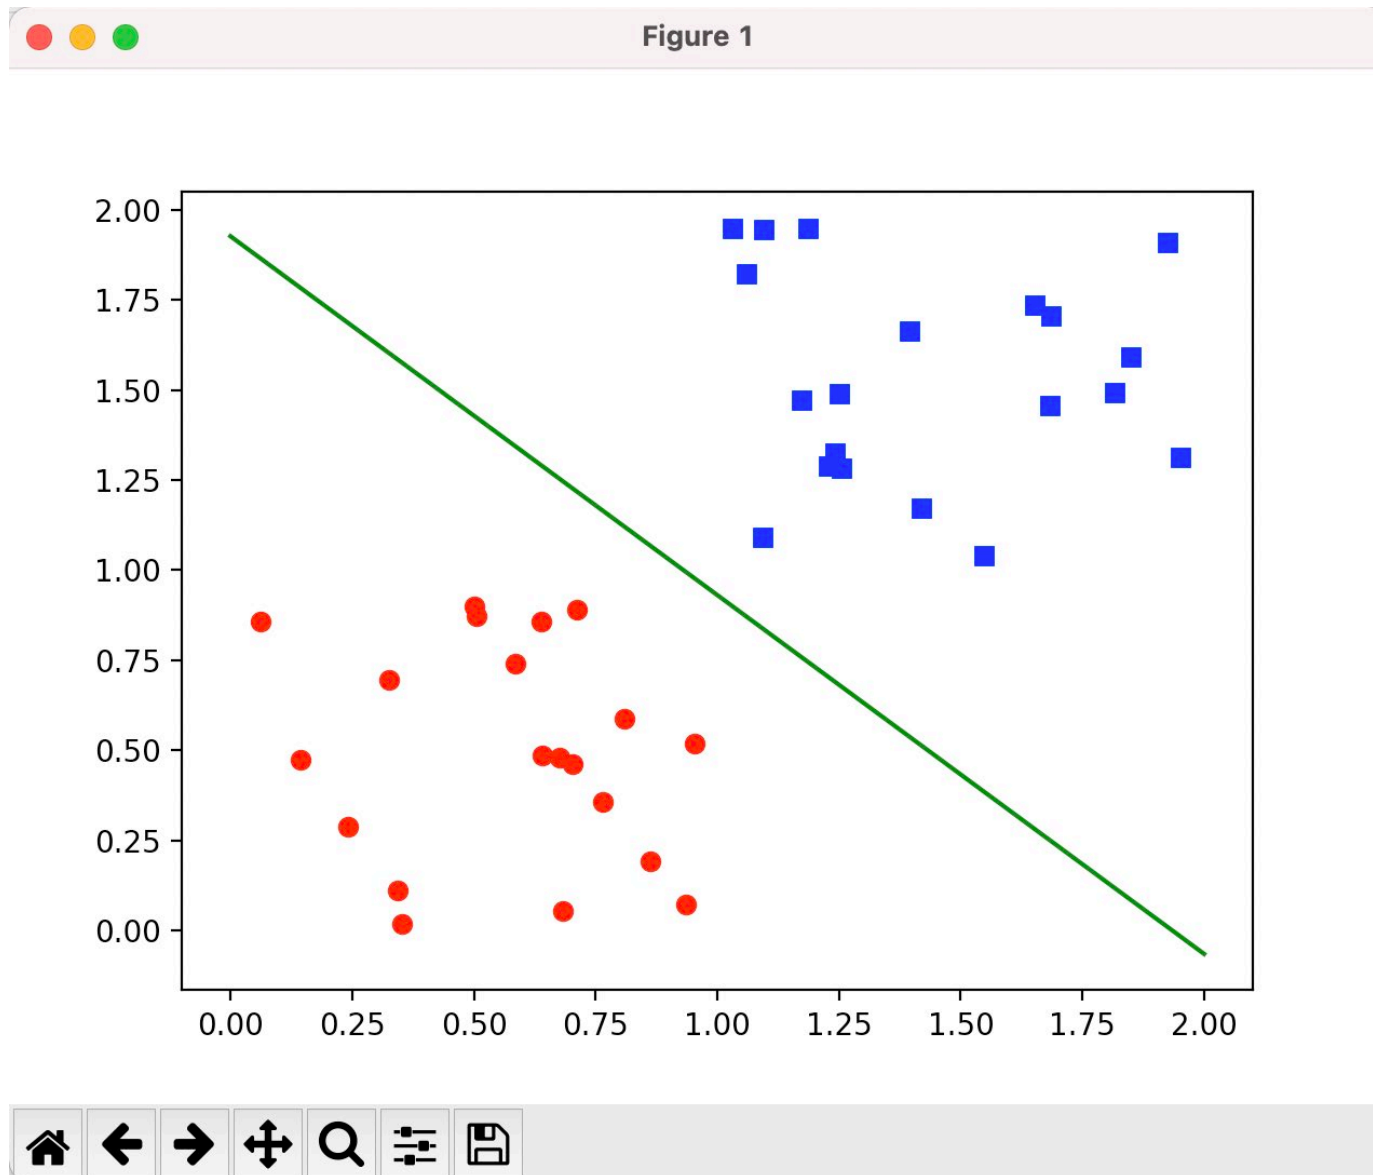
```
plt.plot(np.array([0,2]), np.array([0,2])*a+b, "g-")
```

```
plt.show()
```


Results

```
loss: 0.6918238344332963
loss: 0.12193492747304553
loss: 0.08102355257708946
loss: 0.0644580154145817
loss: 0.05511618732298669
loss: 0.0489784090151211
loss: 0.0445705016555993
loss: 0.04121452216444617
loss: 0.03855167714157677
loss: 0.036372816541204736
Average classification accuracy: 1.0
```

Results



Exploration

How is the decision boundary changed during the learning process?

What will happen if the data are not linearly separable?

What is the effect of using different learning rates?