


ECE/CS 559 - Neural Networks.

Generative Models.


①

- * Two main approaches in machine learning: generative vs. discriminative approach.
- * Observation: X , Target: Y
(e.g. an input image) (e.g. a class label).
- * A discriminative model can provide $P(Y/X=x)$
(e.g. given the input, what are the likelihoods of different class labels).
- * Whereas in generative model, we are interested in finding $P(X/Y=y)$
Given label, generate some samples ~~that~~ belonging to that label (kind of an inverse problem).

Discriminative model :



Generative model



* Noise could be considered as a "seed". Different noises would generate different images of 3s.


Applications of generative models

2

* Low ~~to~~ high resolution image synthesis.

* Text to image translation

"A guy running with an apple on his hand"




* Speech synthesis.

* Error correction

;

3
A simple first idea: Autoencoders.



IDEA: Desired output for a given input X

Input X .

The network learns to compress/encode X to a very low-dimensional representation (through the encoder layers) and then reconstruct the original X through the decoder layers.

After training, one can use the decoder part of the network as a "generator."

Generative Adversarial Networks. (4)

Introduced by Goodfellow et al 2014.

* Two main ideas:

A discriminator network D .

A generator network G .

Ideally $D(x) = \begin{cases} 1, & \text{if } x \text{ is a real image} \\ 0, & \text{if } x \text{ is a fake image} \end{cases}$
provided by e.g. the generator network.

In general, $D(x)$ is the likelihood of an image being real.
 G takes noise z as an input and generates $G(z)$, a generated image.

* Define the objective

$$\min_{G} \max_{D} E[\log D(x) + \log(1 - D(G(z)))]$$

* Theorem: The solution satisfies $p_{G(z)}(x) = p_X(x)$.

Fence, the ^{probability density} output of the generator matches the input data distribution.

Proof (Sketch) ① For a fixed G , find the optimal discriminator.

② Optimize over G .

How to train GANs?

⑤.

- ① Generate data samples x_1, \dots, x_m
- ② Generate noise samples z_1, \dots, z_m .
- ③ Update the discriminator by ascending its stochastic gradient.

$$\nabla_{\theta_D} \frac{1}{m} \sum_{i=1}^m (\log D(x_i) + \log(1 - D(G(z_i))))$$

↓ discriminator parameters.

- ④ Generate new noise samples z_1, \dots, z_m
- ⑤ Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m \log(1 - D(G(z_i)))$$

↓ generator parameters

- ⑥ Goto ① until convergence.

At convergence, one arrives at a solution where $p_{G(z)}(x) = p_X(x)$ and $D(x) = \frac{1}{2} \forall x$. (all images are equally likely to be real or fake as the generator is perfect.)

6 Taken from Goodfellow et al.'s paper.

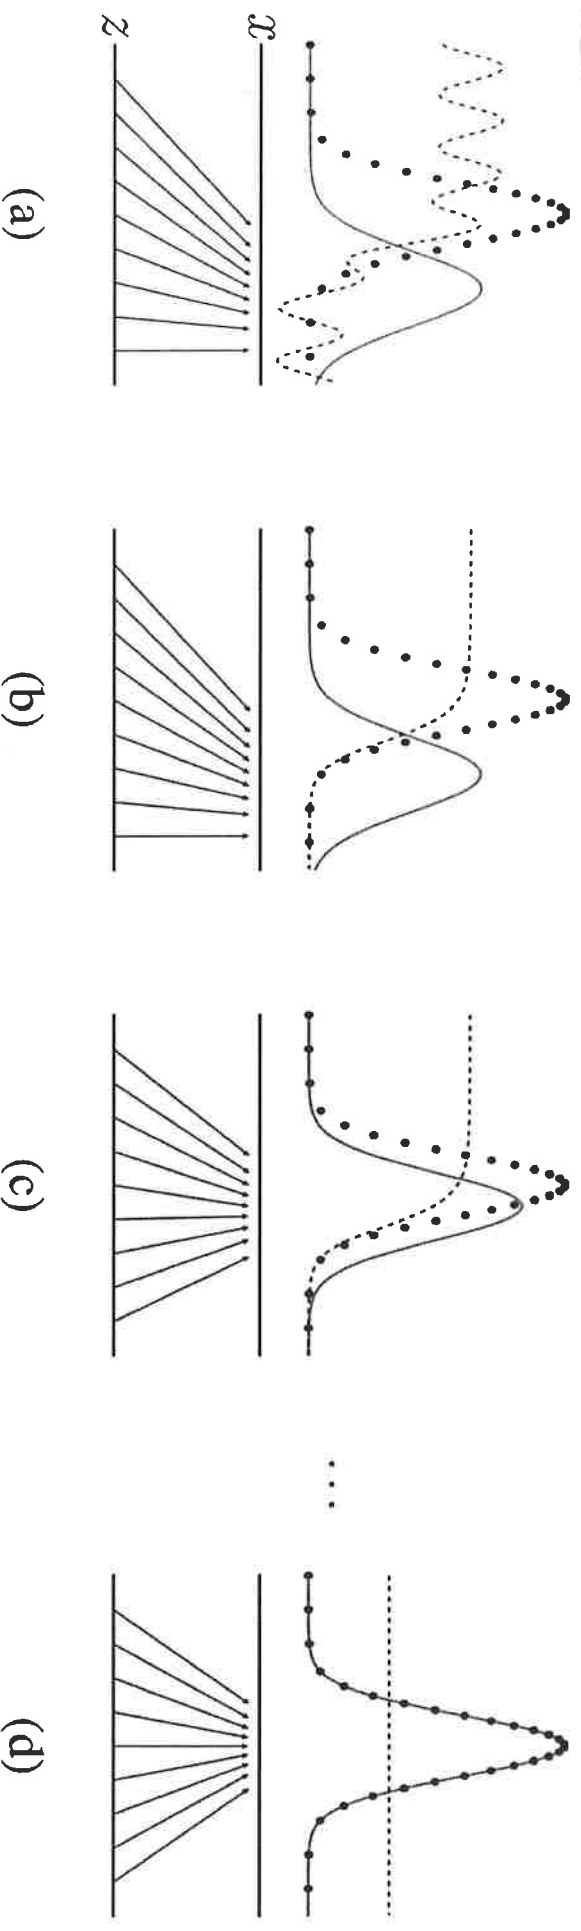


Figure 1: Generative adversarial nets are trained by simultaneously updating the discriminative distribution (D , blue, dashed line) so that it discriminates between samples from the data generating distribution (black, dotted line) p_x from those of the generative distribution p_g (G) (green, solid line). The lower horizontal line is the domain from which z is sampled, in this case uniformly. The horizontal line above is part of the domain of x . The upward arrows show how the mapping $x = G(z)$ imposes the non-uniform distribution p_g on transformed samples. G contracts in regions of high density and expands in regions of low density of p_g . (a) Consider an adversarial pair near convergence: p_g is similar to p_{data} and D is a partially accurate classifier. (b) In the inner loop of the algorithm D is trained to discriminate samples from data, converging to $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}$. (c) After an update to G , gradient of D has guided $G(z)$ to flow to regions that are more likely to be classified as data. (d) After several steps of training, if G and D have enough capacity, they will reach a point at which both cannot improve because $p_g = p_{\text{data}}$. The discriminator is unable to differentiate between the two distributions, i.e. $D(x) = \frac{1}{2}$.

Conditional GANs:

⑦.

GANs cannot natively generate images for a given label. E.g. "Generate an image of a 3". For this purpose, we can use a conditional GAN. All that has to be done is to feed the class label (e.g. as a one-hot-encoded vector) to the generator as well as the discriminator during training (and, of course, during generation). The class label thus becomes an extra input to both D and G .